

Technological Institute of the Philippines Quezon City - Computer Engineering	
Course Code:	CPE 019
Code Title:	Emerging Technologies in CpE 2
2nd Semester	AY 2024 - 2025

ACTIVITY NO. 6.2	Training Neural Networks
Name	Calvadores, Kelly Joseph
Section	CPE32S3
Date Performed:	March 30, 2024
Date Submitted:	April 2, 2024
Instructor:	Engr. Roman M. Richard

```
# Import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_auc_score, roc_curve, accuracy_score
from sklearn.ensemble import RandomForestClassifier
from keras.models import Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
from keras.optimizers import Adam, SGD, RMSprop
%matplotlib inline
```

Load Dataset

```
names = ["times_pregnant", "glucose_tolerance_test", "blood_pressure", "skin_thickness", "insulin",
         "bmi", "pedigree_function", "age", "has_diabetes"]
diabetes_df = pd.read_csv("pima-indians-diabetes.csv", names = names)
```

Check the top 5 samples of the data

```
print(diabetes_df.shape)
diabetes_df.sample(5)
```

(768, 9)

	times_pregnant	glucose_tolerance_test	blood_pressure	skin_thickness	insulin
349	5	0	80	32	0
187	1	128	98	41	58
517	7	125	86	0	0
665	1	112	80	45	132
230	4	142	86	0	0

```
diabetes_df.dtypes

times_pregnant      int64
glucose_tolerance_test  int64
blood_pressure      int64
skin_thickness      int64
insulin             int64
bmi                 float64
pedigree_function   float64
age                 int64
has_diabetes        int64
dtype: object

X = diabetes_df.iloc[:, :-1].values
y = diabetes_df["has_diabetes"].values
```

Split the data to Train, and Test (75%, 25%)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=11111)
```

```
np.mean(y), np.mean(1-y)
```

```
(0.3489583333333333, 0.6510416666666666)
```

✓ Normalize the data

```
normalizer = StandardScaler()
X_train_norm = normalizer.fit_transform(X_train)
X_test_norm = normalizer.transform(X_test)
```

✓ Define the model:

```
model = Sequential([
    Dense(12, input_shape=(8,), activation="relu"),
    Dense(1, activation="sigmoid")
])
```

✓ View the model summary

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	108
dense_1 (Dense)	(None, 1)	13
Total params: 121 (484.00 Byte)		
Trainable params: 121 (484.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

✓ Train the model

```
model.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])
run_hist_1 = model.fit(X_train_norm, y_train, validation_data=(X_test_norm, y_test), epochs=200)
```

```
WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g., tf.keras.optimiz
Epoch 1/200
18/18 [=====] - 1s 14ms/step - loss: 0.7711 - accuracy: 0.4167 - val_loss: 0.7685 - val_accuracy: 0.4427
Epoch 2/200
18/18 [=====] - 0s 4ms/step - loss: 0.7405 - accuracy: 0.4497 - val_loss: 0.7423 - val_accuracy: 0.4583
Epoch 3/200
18/18 [=====] - 0s 4ms/step - loss: 0.7159 - accuracy: 0.4948 - val_loss: 0.7210 - val_accuracy: 0.4896
Epoch 4/200
18/18 [=====] - 0s 4ms/step - loss: 0.6956 - accuracy: 0.5469 - val_loss: 0.7030 - val_accuracy: 0.5938
Epoch 5/200
18/18 [=====] - 0s 4ms/step - loss: 0.6785 - accuracy: 0.5990 - val_loss: 0.6876 - val_accuracy: 0.6198
Epoch 6/200
18/18 [=====] - 0s 3ms/step - loss: 0.6638 - accuracy: 0.6424 - val_loss: 0.6742 - val_accuracy: 0.6615
Epoch 7/200
18/18 [=====] - 0s 4ms/step - loss: 0.6511 - accuracy: 0.6858 - val_loss: 0.6623 - val_accuracy: 0.6875
Epoch 8/200
18/18 [=====] - 0s 4ms/step - loss: 0.6403 - accuracy: 0.7118 - val_loss: 0.6519 - val_accuracy: 0.6875
Epoch 9/200
18/18 [=====] - 0s 4ms/step - loss: 0.6305 - accuracy: 0.7222 - val_loss: 0.6423 - val_accuracy: 0.6823
Epoch 10/200
18/18 [=====] - 0s 4ms/step - loss: 0.6216 - accuracy: 0.7257 - val_loss: 0.6336 - val_accuracy: 0.7031
Epoch 11/200
18/18 [=====] - 0s 4ms/step - loss: 0.6136 - accuracy: 0.7344 - val_loss: 0.6256 - val_accuracy: 0.7188
Epoch 12/200
18/18 [=====] - 0s 4ms/step - loss: 0.6062 - accuracy: 0.7448 - val_loss: 0.6182 - val_accuracy: 0.7292
Epoch 13/200
18/18 [=====] - 0s 4ms/step - loss: 0.5995 - accuracy: 0.7535 - val_loss: 0.6114 - val_accuracy: 0.7396
Epoch 14/200
18/18 [=====] - 0s 4ms/step - loss: 0.5932 - accuracy: 0.7483 - val_loss: 0.6051 - val_accuracy: 0.7292
Epoch 15/200
18/18 [=====] - 0s 4ms/step - loss: 0.5873 - accuracy: 0.7535 - val_loss: 0.5993 - val_accuracy: 0.7240
```

```

Epoch 16/200
18/18 [=====] - 0s 4ms/step - loss: 0.5818 - accuracy: 0.7517 - val_loss: 0.5938 - val_accuracy: 0.7292
Epoch 17/200
18/18 [=====] - 0s 4ms/step - loss: 0.5766 - accuracy: 0.7517 - val_loss: 0.5887 - val_accuracy: 0.7292
Epoch 18/200
18/18 [=====] - 0s 4ms/step - loss: 0.5717 - accuracy: 0.7535 - val_loss: 0.5839 - val_accuracy: 0.7292
Epoch 19/200
18/18 [=====] - 0s 4ms/step - loss: 0.5671 - accuracy: 0.7535 - val_loss: 0.5793 - val_accuracy: 0.7292
Epoch 20/200
18/18 [=====] - 0s 4ms/step - loss: 0.5627 - accuracy: 0.7535 - val_loss: 0.5751 - val_accuracy: 0.7292
Epoch 21/200
18/18 [=====] - 0s 3ms/step - loss: 0.5586 - accuracy: 0.7517 - val_loss: 0.5711 - val_accuracy: 0.7292
Epoch 22/200
18/18 [=====] - 0s 4ms/step - loss: 0.5547 - accuracy: 0.7535 - val_loss: 0.5674 - val_accuracy: 0.7292
Epoch 23/200
18/18 [=====] - 0s 3ms/step - loss: 0.5511 - accuracy: 0.7535 - val_loss: 0.5638 - val_accuracy: 0.7344
Epoch 24/200
18/18 [=====] - 0s 3ms/step - loss: 0.5476 - accuracy: 0.7535 - val_loss: 0.5604 - val_accuracy: 0.7344
Epoch 25/200
18/18 [=====] - 0s 4ms/step - loss: 0.5442 - accuracy: 0.7535 - val_loss: 0.5572 - val_accuracy: 0.7344
Epoch 26/200
18/18 [=====] - 0s 3ms/step - loss: 0.5412 - accuracy: 0.7552 - val_loss: 0.5541 - val_accuracy: 0.7396
Epoch 27/200
18/18 [=====] - 0s 4ms/step - loss: 0.5381 - accuracy: 0.7535 - val_loss: 0.5512 - val_accuracy: 0.7396
Epoch 28/200
18/18 [=====] - 0s 4ms/step - loss: 0.5351 - accuracy: 0.7535 - val_loss: 0.5483 - val_accuracy: 0.7396

```

```

y_pred_class_nn_1 = np.argmax(model.predict(X_test_norm), axis=1)
y_pred_prob_nn_1 = model.predict(X_test_norm)

```

```

6/6 [=====] - 0s 3ms/step
6/6 [=====] - 0s 3ms/step

```

```

y_pred_class_nn_1[:10]

array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

```

y_pred_prob_nn_1[:10]

array([[0.57572615],
       [0.66893107],
       [0.2787839 ],
       [0.15336382],
       [0.21870871],
       [0.52485234],
       [0.02949159],
       [0.19661158],
       [0.9426653 ],
       [0.11621345]], dtype=float32)

```

✓ Create the plot_roc function

```

def plot_roc(y_test, y_pred, model_name):
    fpr, tpr, thr = roc_curve(y_test, y_pred)
    fig, ax = plt.subplots(figsize=(8, 8))
    ax.plot(fpr, tpr, 'k-')
    ax.plot([0, 1], [0, 1], 'k--', linewidth=.5) # roc curve for random model
    ax.grid(True)
    ax.set(title='ROC Curve for {} on PIMA diabetes problem'.format(model_name),
           xlim=[-0.01, 1.01], ylim=[-0.01, 1.01])

```

✓ Evaluate the model performance and plot the ROC CURVE


```

print('accuracy is {:.3f}'.format(accuracy_score(y_test,y_pred_class_nn_1)))
print('roc_auc is {:.3f}'.format(roc_auc_score(y_test,y_pred_prob_nn_1)))

plot_roc(y_test, y_pred_prob_nn_1, 'NN')

```

accuracy is 0.641
roc-auc is 0.822




Plot the training loss and the validation loss over the different epochs and see how it looks

```
run_hist_1.history.keys()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
fig, ax = plt.subplots()
ax.plot(run_hist_1.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(run_hist_1.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

```
<matplotlib.legend.Legend at 0x7c3fc0c16410>
```



What is your interpretation about the result of the train and validation loss?

- The result of the train and validation loss are having big gaps by the end of the graph, this means that the data is overfitting due to noises and outliers that may impact the dataset.

✓ Supplementary Activity

```
# Import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, precision_recall_curve, roc_auc_score, roc_curve, accuracy_score
from sklearn.ensemble import RandomForestClassifier
from keras.models import Sequential
from keras.layers import Input, Dense, Flatten, Dropout, BatchNormalization
from keras.optimizers import Adam, SGD, RMSprop
%matplotlib inline

diabetes = pd.read_csv("pima-indians-diabetes.csv", names = names)
names = ["times_pregnant", "glucose_tolerance_test", "blood_pressure", "skin_thickness", "insulin",
        "bmi", "pedigree_function", "age", "has_diabetes"]

diabetes.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   times_pregnant         768 non-null    int64
1   glucose_tolerance_test 768 non-null    int64
2   blood_pressure         768 non-null    int64
3   skin_thickness         768 non-null    int64
4   insulin                768 non-null    int64
5   bmi                   768 non-null    float64
6   pedigree_function      768 non-null    float64
7   age                   768 non-null    int64
8   has_diabetes           768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

X = diabetes_df.iloc[:, :-1].values
y = diabetes_df["has_diabetes"].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5, random_state=11111)

normalize = StandardScaler()
XTrainNorm = normalize.fit_transform(X_train)
XTestNorm = normalize.transform(X_test)

#Build a model with two hidden layers, each with 6 nodes
#Use the "relu" activation function for the hidden layers, and "sigmoid" for the final layer

model = Sequential([
    Dense(6, input_shape=(8,)), activation="relu"),
    Dense(1, activation="sigmoid")
])

model.summary()

Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 6)	54
dense_6 (Dense)	(None, 1)	7

```

=====
Total params: 61 (244.00 Byte)
Trainable params: 61 (244.00 Byte)
Non-trainable params: 0 (0.00 Byte)
=====

#Use a learning rate of .003 and train for 1500 epochs
model.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])
ModelSupp = model.fit(XTrainNorm, y_train, validation_data=(XTestNorm, y_test), epochs=1500)

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-1-8006da3ef532> in <cell line: 2>()
      1 #Use a learning rate of .003 and train for 1500 epochs
----> 2 model.compile(SGD(lr = .003), "binary_crossentropy", metrics=["accuracy"])
      3 ModelSupp = model.fit(XTrainNorm, y_train, validation_data=(XTestNorm,
y_test), epochs=1500)

NameError: name 'model' is not defined

```

```

y_pred_class_nn_1 = np.argmax(model.predict(XTestNorm), axis=1)
y_pred_prob_nn_1 = model.predict(XTestNorm)

```

```

12/12 [=====] - 0s 2ms/step
12/12 [=====] - 0s 2ms/step

```

```

def plot_roc(y_test, y_pred, model_name):
    fpr, tpr, thr = roc_curve(y_test, y_pred)
    fig, ax = plt.subplots(figsize=(8, 8))
    ax.plot(fpr, tpr, 'k-')
    ax.plot([0, 1], [0, 1], 'k--', linewidth=1.5) # roc curve for random model
    ax.grid(True)
    ax.set(title='ROC Curve for {} on PIMA diabetes problem'.format(model_name),
           xlim=[-0.01, 1.01], ylim=[-0.01, 1.01])

```

```

#Plot the roc curve for the predictions
print('accuracy is {:.3f}'.format(accuracy_score(y_test,y_pred_class_nn_1)))
print('roc-auc is {:.3f}'.format(roc_auc_score(y_test,y_pred_prob_nn_1)))


```

```
plot_roc(y_test, y_pred_prob_nn_1, 'NN')
```

```

accuracy is 0.667
roc-auc is 0.809

```




```
ModelSupp.history.keys()
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
#Graph the trajectory of the loss functions, accuracy on both train and test set
fig, ax = plt.subplots()
ax.plot(ModelSupp.history["loss"], 'r', marker='.', label="Train Loss")
ax.plot(ModelSupp.history["val_loss"], 'b', marker='.', label="Validation Loss")
ax.legend()
```

```
<matplotlib.legend.Legend at 0x7c3fc0696cb0>
```



Remarks

In this graph, it has become overfitting but this time the dataset's train loss and validation loss create a huge gap to its graph by the end of it. The Accuracy of this data is no more than 0.677 and roc is much higher, which indicates that the data get the noise and and did not perform well.

✓ Conclusion

In this activity, i work on demonstrating the training of neural networks using keras. I been able to train neural networks by building models using keras library, this difficult to train models you to some adjusting values such as weights learning rates and other codes, but are you able to adapt and understand training neural network in the given time but in some circumstances the models needs more time to able the data be more accurate.